



**המחלקה להנדסת חשמל וALKتروניקה**

**(31245) מערכות למדות ולמידה عمוקה**

**Lab 7 report**


---

**פרנסיס עבוד**

**Prepared for: Dr. Amer Adler**

**Date: 24/05/2025**

In this lab we will implement a CNN for classifying 10 objects classes from CIFAR10 dataset, each image in the datasets is 32 X 32 pixels in RGB:



```

1 # Lab Question 1 & 2: Setup and Load Data
2 # 1. Import libraries
3 import tensorflow as tf
4 from tensorflow.keras import datasets, layers, models
5 import matplotlib.pyplot as plt
6 import numpy as np
7
8 # 2. Load and preprocess CIFAR-10 data
9 (x_train, y_train), (x_test, y_test) = datasets.cifar10.load_data()
10 x_train, x_test = x_train / 255.0, x_test / 255.0 # Normalize pixel values to be between 0 and 1
11
12 # Lab Question 2: Display sample images from CIFAR-10
13 class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
14                 'dog', 'frog', 'horse', 'ship', 'truck']
15
16 plt.figure(figsize=(10,10)) # Adjusted figure size for better layout
17 for i in range(25): # Display the first 25 images
18     plt.subplot(5,5,i+1)
19     plt.xticks([])
20     plt.yticks([])
21     plt.grid(False)
22     plt.imshow(x_train[i])
23     # The CIFAR labels happen to be arrays, so you need the first element.
24     plt.xlabel(class_names[y_train[i][0]])
25 plt.tight_layout() # Adjusts subplot params for a tight layout.
--  plt.show()

```



### Ex.3:

The network model is given by:

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))
```

### Code:

```
27 # Lab Question 3: Define the CNN model structure
28 # The model structure described in Question 3 is implemented in this function.
29
30 # Original model from question: Conv2D(32,(3,3)) -> MaxPool -> Conv2D(64,(3,3)) -> MaxPool -> Conv2D(64,(3,3)) -> Flatten -> Dense(64) -> Dense(10)
31 v def build_model(conv_layer_specs, first_kernel_size=(3,3), dense_units=64):
32     model = models.Sequential()
33     # First Conv Layer (as per Question 3)
34     model.add(layers.Conv2D(32, first_kernel_size, activation='relu', input_shape=(32, 32, 3)))
35     model.add(layers.MaxPooling2D((2, 2)))
36
37     # Subsequent Convolutional Layers based on specs
38     # Each spec can be a tuple: (filters, add_pooling_after)
39     # For the original model, conv_layer_specs would be: [(64, True), (64, False)]
40     # True means add MaxPooling2D after this Conv2D layer
41     # False means do not add MaxPooling2D after this Conv2D layer (e.g., before Flatten)
42     v for i, spec in enumerate(conv_layer_specs):
43         filters_count, add_pooling_after = spec
44         model.add(layers.Conv2D(filters_count, (3, 3), activation='relu'))
45         if add_pooling_after:
46             |   model.add(layers.MaxPooling2D((2, 2)))
47
48     model.add(layers.Flatten())
49     model.add(layers.Dense(dense_units, activation='relu'))
50     model.add(layers.Dense(10)) # Output layer
51     return model
```

### Ex.4:

Train the given network, using stochastic gradient descent with 20 epochs:

- a. Replace the optimizer in the method “compile()” to ‘sgd’.
- b. In the method “fit()” Change the validation set to be 2.5% of the training data, using “validation\_split = 0.025”, and set the number of epochs to 20.
- c. Plot a graph of the training & validation accuracy vs. number of epochs.
- d. Report the testing set accuracy.


## Code:

```
53     # Lab Question 4: Train the original network and evaluate
54     # This function handles training and evaluation, including requirements from Question 4.
55     def train_and_evaluate(conv_layer_specs_config, first_kernel_size_config=(3,3), dense_units_config=64, epochs_config=20, title_prefix='Model'):
56         model = build_model(conv_layer_specs_config, first_kernel_size_config, dense_units_config)
57
58         # Lab Question 4.a: Replace the optimizer in the method "compile()" to 'sgd'.
59         model.compile(optimizer='sgd',
60                       loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
61                       metrics=['accuracy'])
62
63         # Lab Question 4.b: In the method "fit()" Change the validation set to be 2.5%
64         # of the training data, using "validation_split = 0.025", and set the
65         # number of epochs to 20.
66         history = model.fit(x_train, y_train, epochs=epochs_config,
67                             validation_split=0.025, batch_size=64, verbose=2)
68
69         test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
70
71         # Lab Question 4.c: Plot a graph of the training & validation accuracy vs. number of epochs.
72         plt.figure()
73         plt.plot(history.history['accuracy'], label='train acc')
74         plt.plot(history.history['val_accuracy'], label='val acc')
75         plt.title(f'{title_prefix} - Accuracy vs. Epochs')
76         plt.xlabel('Epoch')
77         plt.ylabel('Accuracy')
78         plt.legend()
79         plt.show()
80
81         # Lab Question 4.d: Report the testing set accuracy.
82         print(f'{title_prefix} - Test accuracy: {test_acc:.4f}')
83         return test_acc
84
85 # 4. Original model (Corresponds to Lab Question 3 & 4)
86 # Training the network as specified in Question 3, with training parameters from Question 4.
87 print("Original Model (Lab Question 3 & 4):")
88 # conv_layer_specs_config:
89 # First 64-filter layer is followed by pooling.
90 # Second 64-filter layer (the 3rd conv layer overall) is NOT followed by pooling before Flatten.
91 original_model_conv_specs = [(64, True), (64, False)]
92 train_and_evaluate(conv_layer_specs_config=original_model_conv_specs, first_kernel_size_config=(3,3), dense_units_config=64, epochs_config=20,
93                     title_prefix='Original Model')
```

## Output:

Test accuracy= 0.6653

```
762/762 - 20s - 27ms/step - accuracy: 0.6760 - loss: 0.9331 - val_accuracy: 0.6536 - val_loss: 0.9999
Epoch 20/20
762/762 - 18s - 23ms/step - accuracy: 0.6829 - loss: 0.9126 - val_accuracy: 0.6960 - val_loss: 0.9429
Original Model - Test accuracy: 0.6653
```



### Ex.5:

Evaluate the impact on CIFAR-10 classification accuracy (measured on the testing set) of the following modifications to the network (you must re-define the CNN for each modification, otherwise, learned weights are kept). For each modification Plot a graph of the training & validation accuracy vs. number of epochs.

- a. Removal of the second convolutional layer
- b. Removal of the third convolutional layer
- c. Removal of the second and third convolutional layer
- d. Increase and decrease the kernel size of the first CNN layer.
- e. Increase and decrease the kernel size of the first convolutional layer.
- f. Increase and decrease of the number of perceptrons in the Dense layer (currently 64).


### A.Code:

```
94     # 5. Modifications (Corresponds to Lab Question 5)
95     print("\nEvaluating Modifications (Lab Question 5):")
96
97     # Lab Question 5.a: Removal of the second convolutional layer
98     # Original specified: Conv(32)-P -> Conv(64)-P -> Conv(64) -> F -> D(64) -> D(10)
99     # "Second convolutional layer" is the first Conv(64). Removing it means:
100    # Conv(32)-P -> Conv(64) -> F -> D(64) -> D(10)
101    # This means one 64-filter layer, not followed by pooling.
102    q5a_conv_specs = [(64, False)]
103    print("Modification 5.a: Remove 2nd Conv Layer (1st 64-filter layer)")
104    train_and_evaluate(conv_layer_specs_config=q5a_conv_specs, first_kernel_size_config=(3,3), dense_units_config=64, title_prefix='Q5.a - Remove 2nd Conv Layer')
```

### A.Output:

```
762/762 - 32s - 42ms/step - accuracy: 0.7463 - loss: 0.7371 - val_accuracy: 0.6608 - val_loss: 0.9494
Q5.a - Remove 2nd Conv Layer - Test accuracy: 0.6683
```

Test accuracy = 0.6683



## B.Code:


```

106     # Lab Question 5.b: Removal of the third convolutional layer
107     # Original specified: Conv(32)-P -> Conv(64)-P -> Conv(64) -> F -> D(64) -> D(10)
108     # "Third convolutional layer" is the second Conv(64). Removing it means:
109     # Conv(32)-P -> Conv(64)-P -> F -> D(64) -> D(10)
110     # This means one 64-filter layer, followed by pooling.
111     q5b_conv_specs = [(64, True)]
112     print("Modification 5.b: Remove 3rd Conv Layer (2nd 64-filter layer)")
113     train_and_evaluate(conv_layer_specs_config=q5b_conv_specs, first_kernel_size_config=(3,3), dense_units_config=64, title_prefix='Q5.b - Remove 3rd Conv Layer')

```

## B.Output:

762/762 - 20s - 27ms/step - accuracy: 0.6829 - loss: 0.9152 - val\_accuracy: 0.6584 - val\_loss: 1.0059  
 Q5.b - Remove 3rd Conv Layer - Test accuracy: 0.6383



## C.Code:


```

115     # Lab Question 5.c: Removal of the second and third convolutional layers
116     # This means removing both 64-filter layers.
117     # Model: Conv(32)-P -> F -> D(64) -> D(10)
118     q5c_conv_specs = [] # No additional conv layers
119     print("Modification 5.c: Remove 2nd & 3rd Conv Layers (both 64-filter layers)")
120     train_and_evaluate(conv_layer_specs_config=q5c_conv_specs, first_kernel_size_config=(3,3), dense_units_config=64, title_prefix='Q5.c - Only 1st Conv Layer')
121

```

## C.Output:

762/762 - 13s - 17ms/step - accuracy: 0.6675 - loss: 0.9635 - val\_accuracy: 0.6400 - val\_loss: 1.0755  
 Q5.c - Only 1st Conv Layer - Test accuracy: 0.6183




## D. and E. Code:

```
122     # Lab Question 5.d & 5.e: Increase and decrease the kernel size of the first CNN layer.  
123     # The "first CNN layer" is the Conv2D(32, kernel_size, ...) layer.  
124     # The rest of the original model structure (including pooling decisions) remains.  
125     print("Modification 5.d/e: Increase kernel size of first conv layer to (5x5)")  
126     train_and_evaluate(conv_layer_specs_config=original_model_conv_specs, first_kernel_size_config=(5,5), dense_units_config=64, title_prefix='Q5.d/e -  
Kernel 5x5 (1st Layer)')  
127  
128     print("Modification 5.d/e: Decrease kernel size of first conv layer to (2x2)")  
129     train_and_evaluate(conv_layer_specs_config=original_model_conv_specs, first_kernel_size_config=(2,2), dense_units_config=64, title_prefix='Q5.d/e -  
Kernel 2x2 (1st Layer)')
```


## D.Output:

```
762/762 - 17s - 23ms/step - accuracy: 0.6738 - loss: 0.9363 - val_accuracy: 0.6576 - val_loss: 1.0068  
Q5.d/e - Kernel 5x5 (1st Layer) - Test accuracy: 0.6394
```



## E.Output:

```
762/762 - 41s - 54ms/step - accuracy: 0.6631 - loss: 0.9654 - val_accuracy: 0.6608 - val_loss: 0.9888  
Q5.d/e - Kernel 2x2 (1st Layer) - Test accuracy: 0.6465
```




## F.Code:

```
131 # Lab Question 5.f: Increase and decrease of the number of perceptrons in the Dense layer (currently 64).  
132 # The rest of the original model structure (including pooling decisions) remains.  
133 print("Modification 5.f: Increase Dense Units to 128")  
134 train_and_evaluate(conv_layer_specs_config=original_model_conv_specs, first_kernel_size_config=(3,3), dense_units_config=128, title_prefix='Q5.f - Dense  
135 128')  
136  
137 print("Modification 5.f: Decrease Dense Units to 32")  
138 train_and_evaluate(conv_layer_specs_config=original_model_conv_specs, first_kernel_size_config=(3,3), dense_units_config=32, title_prefix='Q5.f - Dense  
32')
```

## F.Output Increase:

Q5.f - Dense 128 - Test accuracy: 0.6071



## F.Output Decrease:

Q5.f - Dense 32 - Test accuracy: 0.6325

