



**המחלקה להנדסת חשמל ואלקטרוניקה**

**(31245) מערכות לומדות ולמידה عمוקה**

## **Lab 8 report**

---

**פרנסיס עבוד**

**Prepared for: Dr. Amer Adler**

**Date: 13/06/2025**

## EX.1:

Load the following example into Colab: “Transfer learning with a pretrained ConvNet”: [https://www.tensorflow.org/tutorials/images/transfer\\_learning](https://www.tensorflow.org/tutorials/images/transfer_learning)

## Code:

```

1  import matplotlib.pyplot as plt
2  import numpy as np
3  import os
4  import tensorflow as tf
5
6  # =====
7  # SECTION 1: Data Loading and Preprocessing (from Tensorflow tutorial)
8  # =====
9
10 # Download and setup dataset
11 _URL = 'https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip'
12 path_to_zip = tf.keras.utils.get_file('cats_and_dogs.zip', origin=_URL, extract=True)
13 PATH = os.path.join(os.path.dirname(path_to_zip), 'cats_and_dogs_extracted', 'cats_and_dogs_filtered')
14 train_dir = os.path.join(PATH, 'train')
15 validation_dir = os.path.join(PATH, 'validation')
16
17 BATCH_SIZE = 32
18 IMG_SIZE = (160, 160)
19
20 # Create datasets
21 train_dataset = tf.keras.utils.image_dataset_from_directory(
22     train_dir,
23     shuffle=True,
24     batch_size=BATCH_SIZE,
25     image_size=IMG_SIZE
26 )
27
28 validation_dataset = tf.keras.utils.image_dataset_from_directory(
29     validation_dir,
30     shuffle=True,
31     batch_size=BATCH_SIZE,
32     image_size=IMG_SIZE
33 )
34
35 class_names = train_dataset.class_names
36
37 # Visualize the data
38 plt.figure(figsize=(10, 10))
39 for images, labels in train_dataset.take(1):
40     for i in range(9):
41         ax = plt.subplot(3, 3, i + 1)
42         plt.imshow(images[i].numpy().astype("uint8"))
43         plt.title(class_names[labels[i]])
44         plt.axis("off")
45 plt.savefig('images/dataset_samples.png', dpi=150, bbox_inches='tight')
46 plt.close() # Close the figure to free memory
47 print("Dataset samples saved to images/dataset_samples.png")
48
49 # Split validation dataset into validation and test
50 val_batches = tf.data.experimental.cardinality(validation_dataset)
51 test_dataset = validation_dataset.take(val_batches // 5)
52 validation_dataset = validation_dataset.skip(val_batches // 5)
53
54 print('Number of validation batches: %d' % tf.data.experimental.cardinality(validation_dataset))
55 print('Number of test batches: %d' % tf.data.experimental.cardinality(test_dataset))
56
57 # Configure dataset for performance
58 AUTOTUNE = tf.data.AUTOTUNE
59 train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
60 validation_dataset = validation_dataset.prefetch(buffer_size=AUTOTUNE)
61 test_dataset = test_dataset.prefetch(buffer_size=AUTOTUNE)
62
63 # Data augmentation
64 data_augmentation = tf.keras.Sequential([
65     tf.keras.layers.RandomFlip('horizontal'),
66     tf.keras.layers.RandomRotation(0.2),
67 ])
68
69 # Visualize augmented images
70 for image, _ in train_dataset.take(1):
71     plt.figure(figsize=(10, 10))
72     first_image = image[0]
73     for i in range(9):
74         ax = plt.subplot(3, 3, i + 1)
75         augmented_image = data_augmentation(tf.expand_dims(first_image, 0))
76         plt.imshow(augmented_image[0] / 255)
77         plt.axis('off')
78 plt.savefig('images/data_augmentation.png', dpi=150, bbox_inches='tight')
79 plt.close() # Close the figure to free memory
80 print("Data augmentation visualization saved to images/data_augmentation.png")
81
82 # Setup preprocessing and base model
83 preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input
84
85 # Create the base model from the pre-trained model MobileNetV2
86 IMG_SHAPE = IMG_SIZE + (3,)
87 base_model = tf.keras.applications.MobileNetV2(input_shape=IMG_SHAPE,
88                                                 include_top=False,
89                                                 weights='imagenet')
90
91 # Freeze the base model (154 layers)
92 base_model.trainable = False

```

```

93
94     # Add classification layers
95     global_average_layer = tf.keras.layers.GlobalAveragePooling2D()
96     prediction_layer = tf.keras.layers.Dense(1)
97
98     # Build the complete model
99     inputs = tf.keras.Input(shape=(160, 160, 3))
100    x = data_augmentation(inputs)
101    x = preprocess_input(x)
102    x = base_model(x, training=False)
103    x = global_average_layer(x)
104    x = tf.keras.layers.Dropout(0.2)(x)
105    outputs = prediction_layer(x)
106    model = tf.keras.Model(inputs, outputs)
107
108    print(f"Total layers in base model: {len(base_model.layers)}")

```

## Output:

### ➤ Dataset sample:




### ➤ Data augmentation:



### Ex.2:

Freeze the base CNN layers (MobileNetV2 with 154 layers), and train only the top layers. Provide a loss graph and accuracy graph (vs. epoch number) for 3 different learning rates: 0.01, 0.001, 0.0001 and using SGD, AdaGrad, Adam and RMSprop optimizers. Which learning rate & optimizer provided the best test-set accuracy results?



## Code:

```

11 // =====
12 // SECTION 2: Feature Extraction with Different Optimizers and Learning Rates
13 // Freeze base CNN layers and train only top layers
14 // Test: SGD, AdaGrad, Adam, RMSprop with learning rates 0.01, 0.001, 0.0001
15 // =====
16
17 optimizers = {
18     'SGD': tf.keras.optimizers.SGD,
19     'Adagrad': tf.keras.optimizers.Adagrad,
20     'Adam': tf.keras.optimizers.Adam,
21     'RMSprop': tf.keras.optimizers.RMSprop
22 }
23
24 learning_rates = [0.01, 0.001, 0.0001]
25 results = {}
26 best_accuracy = 0
27 best_optimizer = None
28 best_lr = None
29
30 print("=" * 80)
31 print("SECTION 2: FEATURE EXTRACTION (BASE MODEL FROZEN)")
32 print("=" * 80)
33
34 for optimizer_name, optimizer_class in optimizers.items():
35     for LearnRate in learning_rates:
36         print("\nTraining with {optimizer_name} optimizer and learning rate {LearnRate}")
37         print("=" * 60)
38
39         # Compile the model
40         model.compile(optimizer=optimizer_class(learning_rate=LearnRate),
41                         loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
42                         metrics=['accuracy'])
43
44         initial_epochs = 10
45
46         # Evaluate before training
47         loss0, accuracy0 = model.evaluate(validation_dataset, verbose=0)
48         print(f"Initial loss: {loss0:.4f}, Initial accuracy: {accuracy0:.4f}")
49
50         # Train the model
51         history = model.fit(train_dataset,
52                             epochs=15,
53                             validation_data=validation_dataset,
54                             verbose=1)
55
56         # Evaluate on test dataset
57         test_loss, test_accuracy = model.evaluate(test_dataset, verbose=0)
58         print(f"Test accuracy: {test_accuracy:.4f}")
59
60         # Store results
61         key = f"{optimizer_name}_lr_{LearnRate}"
62         results[key] = {
63             'history': history,
64             'test_accuracy': test_accuracy,
65             'optimizer': optimizer_name,
66             'learning_rate': LearnRate
67         }
68
69 # Track best combination
70 if test_accuracy > best_accuracy:
71     best_accuracy = test_accuracy
72     best_optimizer = optimizer_name
73     best_lr = LearnRate
74
75 # Extract training history
76 acc = history.history['accuracy']
77 val_acc = history.history['val_accuracy']
78 loss = history.history['loss']
79 val_loss = history.history['val_loss']
80
81 # Plot training results
82 plt.figure(figsize=(8, 8))
83 plt.subplot(2, 1, 1)
84 plt.plot(acc, label='Training Accuracy')
85 plt.plot(val_acc, label='Validation Accuracy')
86 plt.ylabel('Accuracy')
87 plt.ylim([min(plt.ylim()),1])
88 plt.title(f"Training and Validation Accuracy vs. number of epochs (LR={LearnRate})")
89 plt.legend(loc='lower right')
90
91 plt.subplot(2, 1, 2)
92 plt.plot(loss, label='Training Loss')
93 plt.plot(val_loss, label='Validation Loss')
94 plt.xlabel('Epoch')
95 plt.ylabel('Cross Entropy')
96 plt.ylim([0,10])
97 plt.title(f"Training and Validation Loss vs. number of epochs (LR={LearnRate})")
98 plt.legend()
99
100 # Save the plot with descriptive filename
101 filename = f'images/training_{optimizer_name}_lr_{LearnRate}.png'
102 plt.savefig(filename, dpi=150, bbox_inches='tight')
103 plt.close() # Close the figure to free memory
104 print(f"Training plot saved to {filename}")
105
106 # Print Section 2 results
107 print("\n" + "=" * 80)
108 print("SECTION 2 RESULTS COMPARISON")
109 print("=" * 80)
110 print(f"Optimizer:{<12} {Learning Rate}:{<15} {Test Accuracy}:{<15}")
111 print("=" * 45)
112
113 for key, result in results.items():
114     optimizer = result['optimizer']
115     lr = result['learning_rate']
116     test_acc = result['test_accuracy']
117     print(f"({optimizer}:{<12} {lr}:{<15} {test_acc}:{<15.4f})")
118
119 print(f"\nBest combination: {best_optimizer} with learning rate {best_lr}")
120 print(f"Best test accuracy: {best_accuracy:.4f}")


```

## Output:

### ➤ SGD:


- Learning rate set to 0.01:

Test accuracy: 0.9635




- Learning rate set to 0.001:

Test accuracy: 0.9688



- Learning rate set to 0.0001:


Test accuracy: 0.9740



➤ AdaGrad:


- Learning rate set to 0.01:

Test accuracy: 0.9792




- Learning rate set to 0.001:

**Test accuracy: 0.9635**



- Learning rate set to 0.0001:


**Test accuracy: 0.9635**



➤ **Adam:**


- Learning rate set to 0.01:

Test accuracy: 0.9688




- Learning rate set to 0.001:

Test accuracy: 0.9740



- Learning rate set to 0.0001:


**Test accuracy: 0.9844**



➤ RMSprop:


- Learning rate set to 0.01:

**Test accuracy: 0.9635**




- Learning rate set to 0.001:

**Test accuracy: 0.9740**



- Learning rate set to 0.0001:

**Test accuracy: 0.9740**



## **RESULTS COMPARISON**

SECTION 2 RESULTS COMPARISON		
Optimizer	Learning Rate	Test Accuracy
SGD	0.01	0.9635
SGD	0.001	0.9688
SGD	0.0001	0.9740
AdaGrad	0.01	0.9792
AdaGrad	0.001	0.9635
AdaGrad	0.0001	0.9635
Adam	0.01	0.9688
Adam	0.001	0.9740
Adam	0.0001	0.9844
RMSprop	0.01	0.9635
RMSprop	0.001	0.9740
RMSprop	0.0001	0.9740

### Ex.3:

Perform fine-tuning of the base network (starting from layer 100, while layers below 100 are frozen) and the inference layers, using the best optimizer from section 2. Use the learning rates from section 2, after division by 10. Which learning rate provided the best test-set accuracy results?

## Code:

```

221  # -----
222  # SECTION 3: Fine-Tuning
223  # Unfreeze from layer 100, keep layers below 100 frozen
224  # Use best optimizer from Section 2
225  # Use learning rates from Section 2 divided by 10: 0.001, 0.0001, 0.00001
226  # -----
227
228  print("\n" + "=" * 88)
229  print("SECTION 3: FINE-TUNING (UNFREEZE FROM LAYER 100)")
230  print("=" * 88)
231
232  # Unfreeze the base model
233  base_model.trainable = True
234
235  # Freeze layers below 100
236  for layer in base_model.layers[:100]:
237      layer.trainable = False
238
239  print(f"Number of layers in the base model: {len(base_model.layers)}")
240  print(f"Fine-tuning from layer 100 onwards")
241
242  # Learning rates for fine-tuning (divided by 10)
243  fine_tune_learning_rates = [lr / 10 for lr in learning_rates] # [0.001, 0.0001, 0.00001]
244
245  fine_tune_results = {}
246  best_fine_tune_accuracy = 0
247  best_fine_tune_lr = None
248
249  # Use the best optimizer from Section 2
250  best_optimizer_class = optimizers[best_optimizer]
251
252  for LearnRate in fine_tune_learning_rates:
253      print(f"\nFine-tuning with {best_optimizer} optimizer and learning rate {LearnRate}")
254      print("-" * 60)
255
256      # Compile the model with fine-tuning learning rate
257      model.compile(optimizer=best_optimizer_class(learning_rate=LearnRate),
258                     loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
259                     metrics=['accuracy'])
260
261      initial_epochs = 10
262
263      # Evaluate before fine-tuning
264      loss0, accuracy0 = model.evaluate(validation_dataset, verbose=0)
265      print(f"Initial loss: {loss0:.4f}, Initial accuracy: {accuracy0:.4f}")
266
267      # Fine-tune the model
268      history = model.fit(train_dataset,
269                           epochs=15,
270                           validation_data=validation_dataset,
271                           verbose=1)
272
273      # Evaluate on test dataset
274      test_loss, test_accuracy = model.evaluate(test_dataset, verbose=0)
275      print(f"Fine-tuned test accuracy: {test_accuracy:.4f}")
276
277
278  # Store fine-tuning results
279  fine_tune_results[LearnRate] = {
280      'history': history,
281      'test_accuracy': test_accuracy
282  }
283
284  # Track best fine-tuning learning rate
285  if test_accuracy > best_fine_tune_accuracy:
286      best_fine_tune_accuracy = test_accuracy
287      best_fine_tune_lr = LearnRate
288
289  # Extract training history
290  acc = history.history['accuracy']
291  val_acc = history.history['val_accuracy']
292  loss = history.history['loss']
293  val_loss = history.history['val_loss']
294
295  # Plot fine-tuning results
296  plt.figure(figsize=(8, 8))
297  plt.subplot(2, 1, 1)
298  plt.plot(acc, label='Training Accuracy')
299  plt.plot(val_acc, label='Validation Accuracy')
300  plt.ylabel('Accuracy')
301  plt.ylim([min(plt.ylim()), 1])
302  plt.title("Training and Validation Accuracy vs. number of epochs (LR={LearnRate})")
303  plt.legend(loc='lower right')
304
305  plt.subplot(2, 1, 2)
306  plt.plot(loss, label='Training Loss')
307  plt.plot(val_loss, label='Validation Loss')
308  plt.xlabel('Epoch')
309  plt.ylabel('Cross Entropy')
310  plt.ylim([0, 1.8])
311  plt.title("Training and Validation Loss vs. number of epochs (LR={LearnRate})")
312  plt.legend()
313
314  # Save the fine-tuning plot
315  filename = f'images/fine_tuning_lr{LearnRate}.png'
316  plt.savefig(filename, dpi=150, bbox_inches='tight')
317  plt.close() # Close the figure to free memory
318  print(f"Fine-tuning plot saved to {filename}")
319
320  # Print Section 3 results
321  print("\n" + "=" * 88)
322  print("SECTION 3 FINE-TUNING RESULTS")
323  print("-" * 88)
324  print(f"{'Learning Rate':<15} {'Test Accuracy':<15}")
325  print("-" * 30)
326
327  for lr, result in fine_tune_results.items():
328      test_acc = result['test_accuracy']
329      print(f"{lr:<15} {test_acc:<15.4f}")
330
331  print("\nBest fine-tuning learning rate: {best_fine_tune_lr}")
332  print("Best fine-tuning test accuracy: {best_fine_tune_accuracy:.4f}")
333  print("Improvement from feature extraction: {best_fine_tune_accuracy - best_accuracy:.4f}")

```


## Output:

```
=====
SECTION 3: FINE-TUNING (UNFREEZE FROM LAYER 100)
=====
Number of layers in the base model: 154
Fine-tuning from layer 100 onwards

Fine-tuning with Adam optimizer and learning rate 0.001
```


- Learning rate set to 0.001:

Fine-tuned test accuracy: 0.9115




- Learning rate set to 0.0001:

Fine-tuned test accuracy: 0.9583



- Learning rate set to 0.00001:

Fine-tuned test accuracy: 0.9688



## FINE-TUNING RESULTS

SECTION 3 FINE-TUNING RESULTS	
Learning Rate	Test Accuracy
0.001	0.9115
0.0001	0.9583
1e-05	0.9688
Best fine-tuning learning rate: 1e-05	
Best fine-tuning test accuracy: 0.9688	
Improvement from feature extraction: -0.0156	

## FINAL SUMMARY

```
FINAL SUMMARY
=====
SECTION 2 (Feature Extraction):
    Best optimizer: Adam
    Best learning rate: 0.0001
    Best test accuracy: 0.9844

SECTION 3 (Fine-tuning):
    Best optimizer: Adam (from Section 2)
    Best test accuracy: 0.9844

SECTION 3 (Fine-tuning):
    Best optimizer: Adam (from Section 2)
SECTION 3 (Fine-tuning):
    Best optimizer: Adam (from Section 2)
    Best learning rate: 1e-05
    Best test accuracy: 0.9688
    Total improvement: -0.0156
```