



**המחלקה להנדסת חשמל ואלקטרוניקה**

**(31245) מערכות לומדות ולמידה عمוקה**

**Lab 10 report**

---

**פרנסיס עבוד**

**Prepared for: Dr. Amer Adler**

**Date: 27/06/2025**

# Deep Learning – Lab – Sentiment Analysis

- 1. Load the MLP-based and 1D-CNN-based IMDB Sentiment Analysis solutions into Colab.**
- 2. Train and compute the test data accuracy. Plot the loss and accuracy training curves.**
- 3. Add early stopping to find the optimal stopping point. Is there any improvement compared to the baseline accuracy?**
- 4. In the CNN network, change the embedding layer dimension to a few different values, re-train and compute the test data accuracy. Which value provides the best accuracy?**
- 5. Repeat step 4, for the MLP network.**
- 6. In the CNN network, remove 1 convolutional layer, and compute accuracy, afterwards remove 2 layers. Try also to change the kernel length. What is the impact on the accuracy?**
- 7. In the MLP network, increase the FC layer to a higher number of neurons, and compute accuracy. Which value provides the best accuracy?**

## Solution:

```
6  ✓ from __future__ import print_function
7  import tensorflow as tf
8  from tensorflow import keras
9  from keras.preprocessing import sequence
10 from keras.models import Sequential
11 from keras.layers import Dense, Embedding, GRU, GlobalAveragePooling1D, Convolution1D, Flatten, Dropout
12 from keras.datasets import imdb
13 from keras.utils import pad_sequences
14 from keras.callbacks import EarlyStopping
15 import matplotlib.pyplot as plt
16 import numpy as np
17 import pandas as pd
18 import os
19
20 # Set random seeds for reproducibility
21 np.random.seed(42)
22 tf.random.set_seed(42)
23
24 # Disable GPU if needed (remove if you want to use GPU)
25 os.environ['CUDA_VISIBLE_DEVICES'] = '0'
26
27 # Global parameters
28 max_features = 10000
29 batch_size = 32
30 max_length = 256
```

```

31
32 print("=*80)
33 print("DEEP LEARNING LAB - SENTIMENT ANALYSIS EXPERIMENTS")
34 print("=*80)
35
36 # Load and prepare data
37 print('\nLoading IMDB data...')
38 (train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=max_features)
39
40 print(f'Training entries: {len(train_data)}, labels: {len(train_labels)})')
41 print(f'Test entries: {len(test_data)}, labels: {len(test_labels)})')
42
43 # Pad sequences
44 train_data = keras.preprocessing.sequence.pad_sequences(train_data, maxlen=max_length, padding='post')
45 test_data = keras.preprocessing.sequence.pad_sequences(test_data, maxlen=max_length, padding='post')
46
47 # Create validation split
48 x_val = train_data[:1000]
49 partial_x_train = train_data[1000:]
50 y_val = train_labels[:1000]
51 partial_y_train = train_labels[1000:]
52
53 print(f'Training set: {len(partial_x_train)}, Validation set: {len(x_val)}, Test set: {len(test_data)})')
54

```

## 1. Load the MLP-based and 1D-CNN-based IMDB Sentiment Analysis solutions into Colab:

### Code:

```

55 # =====
56 # QUESTION 1: Load MLP-based and 1D-CNN-based IMDB Sentiment Analysis solutions
57 # =====
58
59 def create_mlp_model(embedding_dim=16, dense_units=16, vocab_size=10000):
60     """Create MLP model for sentiment analysis"""
61     model = Sequential([
62         keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
63         keras.layers.GlobalAveragePooling1D(),
64         keras.layers.Dense(dense_units, activation='relu'),
65         keras.layers.Dense(1, activation='sigmoid')
66     ])
67     return model
68
69 def create_cnn_model(embedding_dim=300, vocab_size=10000):
70     """Create CNN model for sentiment analysis"""
71     model = Sequential([
72         Embedding(vocab_size, embedding_dim, input_length=max_length),
73         Convolution1D(64, 3, padding='same', activation='relu'),
74         Convolution1D(32, 3, padding='same', activation='relu'),
75         Convolution1D(16, 3, padding='same', activation='relu'),
76         Flatten(),
77         Dropout(0.2),
78         Dense(128, activation='relu'),
79         Dropout(0.2),
80         Dense(1, activation='sigmoid')
81     ])
82     return model
83
84 def train_and_evaluate_model(model, model_name, epochs=40, save_plots=True, use_early_stopping=False):
85     """Train and evaluate a model"""
86     print(f"\n{'='*20} Training {model_name} {'='*20}")
87
88     model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
89     model.summary()
90
91     callbacks = []
92     if use_early_stopping:
93         early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
94         callbacks.append(early_stopping)
95
96     history = model.fit(
97         partial_x_train, partial_y_train,
98         epochs=epochs,
99         batch_size=512,
100        validation_data=(x_val, y_val),
101        verbose=1,
102        callbacks=callbacks
103    )

```

```

105     # Evaluate on test data
106     test_results = model.evaluate(test_data, test_labels, verbose=0)
107     test_accuracy = test_results[1] * 100
108
109     print(f"\n{model_name} Test Accuracy: {test_accuracy:.2f}%")
110
111     if save_plots:
112         # Plot training curves
113         fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 5))
114
115         # Loss plot
116         ax1.plot(history.history['loss'], 'bo-', label='Training loss')
117         ax1.plot(history.history['val_loss'], 'ro-', label='Validation loss')
118         ax1.set_title(f'{model_name} - Training and Validation Loss')
119         ax1.set_xlabel('Epochs')
120         ax1.set_ylabel('Loss')
121         ax1.legend()
122         ax1.grid(True)
123
124         # Accuracy plot
125         ax2.plot(history.history['accuracy'], 'bo-', label='Training accuracy')
126         ax2.plot(history.history['val_accuracy'], 'ro-', label='Validation accuracy')
127         ax2.set_title(f'{model_name} - Training and Validation Accuracy')
128         ax2.set_xlabel('Epochs')
129
130         ax2.set_ylabel('Accuracy')
131         ax2.legend()
132         ax2.grid(True)
133
134         plt.tight_layout()
135         plt.savefig(f'{model_name.lower().replace(" ", "_")}_training_curves.png', dpi=300, bbox_inches='tight')
136         plt.close()
137         print(f"Training curves saved as {model_name.lower().replace(' ', '_')}_training_curves.png")
138
139     return test_accuracy, history

```

## 2. Train and compute the test data accuracy. Plot the loss and accuracy training curves.


### Code:

```

140  # =====
141  # QUESTION 2: Train and compute test data accuracy. Plot loss and accuracy curves.
142  # =====
143  print("\n" + "="*80)
144  print("QUESTION 2: Baseline Model Training and Evaluation")
145  print("="*80)
146
147  # Train baseline MLP model
148  mlp_baseline = create_mlp_model()
149  mlp_baseline_accuracy, mlp_baseline_history = train_and_evaluate_model(
150      mlp_baseline, "Question 2 - MLP Baseline", epochs=40, save_plots=True
151  )
152
153  # Train baseline CNN model
154  cnn_baseline = create_cnn_model()
155  cnn_baseline_accuracy, cnn_baseline_history = train_and_evaluate_model(
156      cnn_baseline, "Question 2 - CNN Baseline", epochs=10, save_plots=True
157  )
158
159  # Save baseline results
160  baseline_results = pd.DataFrame({
161      'Model': ['MLP Baseline', 'CNN Baseline'],
162      'Test Accuracy (%)': [mlp_baseline_accuracy, cnn_baseline_accuracy]
163  })
164  baseline_results.to_csv('question_2_baseline_results.csv', index=False)
165  print(f"\nBaseline results saved to question_2_baseline_results.csv")
166  print(baseline_results)

```

## Output for Q-2:




**Model, Test Accuracy (%)**  
MLP Baseline, 87.52400279045105  
CNN Baseline, 86.98800206184387

### 3. Add early stopping to find the optimal stopping point. Is there any improvement compared to the baseline accuracy?

#### Code:

```
169 # QUESTION 3: Add early stopping to find optimal stopping point
170 # =====
171 print("\n" + "*80")
172 print("QUESTION 3: Early Stopping Implementation")
173 print("*80")
174
175 # Train MLP with early stopping
176 mlp_early_stop = create_mlp_model()
177 mlp_early_accuracy, mlp_early_history = train_and_evaluate_model(
178     mlp_early_stop, "Question 3 - MLP Early Stopping", epochs=40,
179     save_plots=True, use_early_stopping=True
180 )
181
182 # Train CNN with early stopping
183 cnn_early_stop = create_cnn_model()
184 cnn_early_accuracy, cnn_early_history = train_and_evaluate_model([
185     cnn_early_stop, "Question 3 - CNN Early Stopping", epochs=20,
186     save_plots=True, use_early_stopping=True
187 ])
188
189 # Compare with baseline
190 early_stopping_results = pd.DataFrame({
191     'Model': ['MLP Baseline', 'MLP Early Stopping', 'CNN Baseline', 'CNN Early Stopping'],
192     'Test Accuracy (%)': [mlp_baseline_accuracy, mlp_early_accuracy,
193                           cnn_baseline_accuracy, cnn_early_accuracy],
194     'Improvement': [0, mlp_early_accuracy - mlp_baseline_accuracy,
195                      0, cnn_early_accuracy - cnn_baseline_accuracy]
196 })
197 early_stopping_results.to_csv('question_3_early_stopping_results.csv', index=False)
198 print(f"\nEarly stopping results saved to question_3_early_stopping_results.csv")
199 print(early_stopping_results)
```

#### Output for Q-3:



```
Model,Test Accuracy (%),Improvement
MLP Baseline,87.52400279045105,0.0
MLP Early Stopping,88.41599822044373,0.8919954299926758
CNN Baseline,86.98800206184387,0.0
CNN Early Stopping,87.76800036430359,0.7799983024597168
```

**4. In the CNN network, change the embedding layer dimension to a few different values, re-train and compute the test data accuracy. Which value provides the best accuracy?**


### Code:

```

203     print("\n" + "*80")
204     print("QUESTION 4: CNN Embedding Dimension Experiments")
205     print("*80")
206
207     embedding_dims = [50, 100, 200, 300, 400, 500]
208     cnn_embedding_results = []
209
210     for dim in embedding_dims:
211         print(f"\nTesting CNN with embedding dimension: {dim}")
212         model = create_cnn_model(embedding_dim=dim)
213         accuracy, _ = train_and_evaluate_model(
214             model, f"Question 4 - CNN Embedding {dim}D", epochs=10, save_plots=False
215         )
216         cnn_embedding_results.append(accuracy)
217
218     # Create results dataframe and plot
219     cnn_embedding_df = pd.DataFrame({
220         'Embedding_Dimension': embedding_dims,
221         'Test_Accuracy': cnn_embedding_results
222     })
223     cnn_embedding_df.to_csv('question_4_cnn_embedding_results.csv', index=False)
224
225     # Plot results
226     plt.figure(figsize=(10, 6))
227     plt.plot(embedding_dims, cnn_embedding_results, 'bo-', linewidth=2, markersize=8)
228     plt.title('Question 4: CNN Performance vs Embedding Dimension')
229     plt.xlabel('Embedding Dimension')
230     plt.ylabel('Test Accuracy (%)')
231     plt.grid(True, alpha=0.3)
232     plt.xticks(embedding_dims)
233     for i, acc in enumerate(cnn_embedding_results):
234         plt.annotate(f'{acc:.2f}%', (embedding_dims[i], acc),
235                     textcoords="offset points", xytext=(0, 10), ha='center')
236     plt.savefig('question_4_cnn_embedding_comparison.png', dpi=300, bbox_inches='tight')
237     plt.close()
238
239     best_cnn_embedding = embedding_dims[np.argmax(cnn_embedding_results)]
240     best_cnn_accuracy = max(cnn_embedding_results)
241     print(f"\nBest CNN embedding dimension: {best_cnn_embedding} with accuracy: {best_cnn_accuracy:.2f}%")
242     print("Results saved to question 4 cnn embedding results.csv and question 4 cnn embedding comparison.png")

```

### Output for Q-4:



```

Embedding_Dimension,Test_Accuracy
50,86.73
100,86.33
200,84.38
300,86.89
400,86.11
500,86.34


```

## 5. Repeat step 4, for the MLP network.

### Code:

```
248 # =====
249 # QUESTION 5: MLP - Change embedding layer dimension to different values
250 # =====
251
252 print("\n" + "="*80)
253 print("QUESTION 5: MLP Embedding Dimension Experiments")
254 print("="*80)
255
256 mlp_embedding_results = []
257
258 for dim in embedding_dims:
259     print(f"\nTesting MLP with embedding dimension: {dim}")
260     model = create_mlp_model(embedding_dim=dim)
261     accuracy, _ = train_and_evaluate_model(
262         model, f"Question 5 - MLP Embedding {dim}D", epochs=20, save_plots=False
263     )
264     mlp_embedding_results.append(accuracy)
265
266 # Create results dataframe and plot
267 mlp_embedding_df = pd.DataFrame({
268     'Embedding_Dimension': embedding_dims,
269     'Test_Accuracy': mlp_embedding_results
270 })
271 mlp_embedding_df.to_csv('question_5_mlp_embedding_results.csv', index=False)
272
273 # Plot results
274 plt.figure(figsize=(10, 6))
275 plt.plot(embedding_dims, mlp_embedding_results, 'ro-', linewidth=2, markersize=8)
276 plt.title('Question 5: MLP Performance vs Embedding Dimension')
277 plt.xlabel('Embedding Dimension')
278 plt.ylabel('Test Accuracy (%)')
279 plt.grid(True, alpha=0.3)
280 plt.xticks(embedding_dims)
281 for i, acc in enumerate(mlp_embedding_results):
282     plt.annotate(f'{acc:.2f}%', (embedding_dims[i], acc),
283                 textcoords='offset points', xytext=(0,10), ha='center')
284 plt.savefig('question_5_mlp_embedding_comparison.png', dpi=300, bbox_inches='tight')
285 plt.close()
286
287 best_mlp_embedding = embedding_dims[np.argmax(mlp_embedding_results)]
288 best_mlp_accuracy = max(mlp_embedding_results)
289 print(f"\nBest MLP embedding dimension: {best_mlp_embedding} with accuracy: {best_mlp_accuracy:.2f}%")
290 print(f"Results saved to question_5_mlp_embedding_results.csv and question_5_mlp_embedding_comparison.png")
```

### Output for Q-5:



## 6. In the CNN network, remove 1 convolutional layer, and compute accuracy, afterwards remove 2 layers. Try also to change the kernel length. What is the impact on the accuracy?


### Code:

```

292 # QUESTION 6: CNN - Remove convolutional layers and change kernel length
293 # =====
294
295 print("\n" + "="*80)
296 print("QUESTION 6: CNN Architecture Experiments")
297 print("="*80)
298
299
300 def create_cnn_variant(num_conv_layers=3, kernel_size=3, embedding_dim=300):
301     """Create CNN variant with different number of layers and kernel size"""
302     model = Sequential()
303     model.add(Embedding(max_features, embedding_dim, input_length=max_length))
304
305     if num_conv_layers >= 1:
306         model.add(Convolution1D(64, kernel_size, padding='same', activation='relu'))
307     if num_conv_layers >= 2:
308         model.add(Convolution1D(32, kernel_size, padding='same', activation='relu'))
309     if num_conv_layers >= 3:
310         model.add(Convolution1D(16, kernel_size, padding='same', activation='relu'))
311
312     model.add(Flatten())
313     model.add(Dropout(0.2))
314     model.add(Dense(180, activation='sigmoid'))
315     model.add(Dropout(0.2))
316     model.add(Dense(1, activation='sigmoid'))
317
318     return model
319
320 # Test different number of convolutional layers
321 conv_layer_experiments = [
322     (3, 3, "3 Conv Layers (Original)"),
323     (2, 3, "2 Conv Layers"),
324     (1, 3, "1 Conv Layer")
325 ]
326
327 conv_layer_results = []
328 for num_layers, kernel_size, description in conv_layer_experiments:
329     print(f"\nTesting CNN with {description}")
330     model = create_cnn_variant(num_conv_layers=num_layers, kernel_size=kernel_size)
331     accuracy, _ = train_and_evaluate_model(
332         model, f"Question 6 - CNN {description}", epochs=10, save_plots=False
333     )
334     conv_layer_results.append(accuracy)
335
336 # Test different kernel sizes with 3 layers
337 kernel_experiments = [
338     (3, 3, "Kernel Size 3"),
339     (3, 5, "Kernel Size 5"),
340     (3, 7, "Kernel Size 7")
341 ]
342
343 kernel_results = []
344 for num_layers, kernel_size, description in kernel_experiments:
345     print(f"\nTesting CNN with {description}")
346     model = create_cnn_variant(num_conv_layers=num_layers, kernel_size=kernel_size)
347     accuracy, _ = train_and_evaluate_model(
348         model, f"Question 6 - CNN {description}", epochs=10, save_plots=False
349     )
350     kernel_results.append(accuracy)
351
352 # Save results
353 cnn_architecture_df = pd.DataFrame({
354     'Configuration': [desc for _, _, desc in conv_layer_experiments] +
355     [desc for _, _, desc in kernel_experiments],
356     'Test_Accuracy': conv_layer_results + kernel_results
357 })
358 cnn_architecture_df.to_csv('question_6_cnn_architecture_results.csv', index=False)
359
360 # Plot results
361 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))
362
363 # Conv layers plot
364 ax1.bar(range(len(conv_layer_results)), conv_layer_results, color=['blue', 'green', 'red'])
365 ax1.set_title('Question 6: CNN Performance vs Number of Conv Layers')
366 ax1.set_xlabel('Configuration')
367 ax1.set_ylabel('Test Accuracy (%)')
368 ax1.set_xticks(range(len(conv_layer_results)))
369 ax1.set_xticklabels([desc for _, _, desc in conv_layer_experiments], rotation=45)
370 ax1.grid(True, alpha=0.3)
371 for i, acc in enumerate(conv_layer_results):
372     ax1.text(i, acc + 0.5, f'{acc:.2f}%', ha='center')
373
374 # Kernel size plot
375 ax2.bar(range(len(kernel_results)), kernel_results, color=['orange', 'purple', 'brown'])
376 ax2.set_title('Question 6: CNN Performance Vs Kernel Size')
377 ax2.set_xlabel('Configuration')
378 ax2.set_ylabel('Test Accuracy (%)')
379 ax2.set_xticks(range(len(kernel_results)))
380 ax2.set_xticklabels([desc for _, _, desc in kernel_experiments], rotation=45)
381 ax2.grid(True, alpha=0.3)
382 for i, acc in enumerate(kernel_results):
383     ax2.text(i, acc + 0.5, f'{acc:.2f}%', ha='center')
384
385 plt.tight_layout()
386 plt.savefig('question_6_cnn_architecture_comparison.png', dpi=300, bbox_inches='tight')
387 plt.close()
388
389 print("\nCNN Architecture results saved to question_6_cnn_architecture_results.csv and question_6_cnn_architecture_comparison.png")
390 print(cnn_architecture_df)

```

### Output for Q-6:



```

Configuration,Test_Accuracy
3 Conv Layers (Original),86.871999502182
2 Conv Layers,85.93599796295166
1 Conv Layer,87.05199956893921
Kernel Size 3,86.02799773216248
Kernel Size 5,86.6599977016449
Kernel Size 7,86.86800003051758

```

**7. In the MLP network, increase the FC layer to a higher number of neurons, and compute accuracy. Which value provides the best accuracy?**


### Code:

```

392 # =====
393 # QUESTION 7: MLP - Increase FC layer neurons
394 # =====
395
396 print("\n" + "="*80)
397 print("QUESTION 7: MLP Dense Layer Size Experiments")
398 print("-"*80)
399
400 dense_units = [8, 16, 32, 64, 128, 256, 512]
401 mlp_dense_results = []
402
403 for units in dense_units:
404     print(f"\nTesting MLP with {units} dense units")
405     model = create_mlp_model(dense_units=units)
406     accuracy, _ = train_and_evaluate_model(
407         model, f"Question 7 - MLP {units} Units", epochs=20, save_plots=False
408     )
409     mlp_dense_results.append(accuracy)
410
411 # Create results dataframe and plot
412 mlp_dense_df = pd.DataFrame({
413     'Dense_Units': dense_units,
414     'Test_Accuracy': mlp_dense_results
415 })
416 mlp_dense_df.to_csv('question_7_mlp_dense_results.csv', index=False)
417
418 # Plot results
419 plt.figure(figsize=(10, 6))
420 plt.plot(dense_units, mlp_dense_results, 'go-', linewidth=2, markersize=8)
421 plt.title('Question 7: MLP Performance vs Dense Layer Size')
422 plt.xlabel('Number of Dense Units')
423 plt.ylabel('Test Accuracy (%)')
424 plt.grid(True, alpha=0.3)
425 plt.xscale('log')
426 plt.xticks(dense_units, dense_units)
427 for i, acc in enumerate(mlp_dense_results):
428     plt.annotate(f'{acc:.2f}%', (dense_units[i], acc),
429                  textcoords="offset points", xytext=(-10, 10), ha='center')
430 plt.savefig('question_7_mlp_dense_comparison.png', dpi=300, bbox_inches='tight')
431 plt.close()
432
433 best_dense_units = dense_units[np.argmax(mlp_dense_results)]
434 best_dense_accuracy = max(mlp_dense_results)
435 print(f"\nBest MLP dense units: {best_dense_units} with accuracy: {best_dense_accuracy:.2f}%")
436 print("Results saved to question_7_mlp_dense_results.csv and question_7_mlp_dense_comparison.png")


```

### Output for Q-7:



Dense_Units, Test_Accuracy
8, 88.10799717903137
16, 88.46799731254578
32, 87.99600005149841
64, 87.94000148773193
128, 87.56800293922424
256, 87.06799745559692
512, 86.00800037384033

# **SUMMARY OF ALL EXPERIMENTS**



Question,Test_Accuracy,Rank
Q7 - MLP Best Dense (16 units),88.46799731254578,1.0
Q3 - MLP Early Stopping,88.41599822044373,2.0
Q3 - CNN Early Stopping,87.76800036430359,3.0
Q5 - MLP Best Embedding (50D),87.76000142097473,4.0
Q2 - MLP Baseline,87.52400279045105,5.0
Q6 - CNN Best Architecture,87.05199956893921,6.0
Q2 - CNN Baseline,86.98800206184387,7.0
Q4 - CNN Best Embedding (300D),86.89200282096863,8.0